

LabAssignment 5

1. Implement logistic regression using Python/R to perform classification on Social_Network_Ads.csv dataset.
2. Compute Confusion matrix to find TP, FP, TN, FN, Accuracy, Error rate, Precision, Recall on the given dataset.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")

url = "https://raw.githubusercontent.com/DSDBA/main/diabetes.csv"
dataset = pd.read_csv(url)
```

dataset

	Age	BMI	Blood Pressure	Glucose Level	Insulin	Outcome
0	25	22.0	80	95	15	0
1	32	27.8	85	120	40	0
2	47	30.5	90	150	60	1
3	55	31.2	88	110	55	1
4	29	24.7	78	100	20	0
5	60	35.0	95	180	70	1
6	40	28.4	82	125	45	0
7	70	34.6	92	160	65	1
8	50	29.0	89	140	50	1
9	38	23.6	85	105	25	0
10	41	33.1	87	130	50	1
11	45	26.3	90	115	35	0
12	51	32.0	91	145	55	1
13	37	28.5	80	105	30	0
14	60	36.0	89	160	65	1
15	30	24.3	78	98	20	0
16	63	34.1	92	155	60	1
17	35	25.5	86	110	40	0
18	56	29.5	93	150	60	1
19	43	28.2	88	120	45	0
20	64	35.6	91	165	70	1
21	33	27.0	84	130	55	0
22	39	30.4	90	140	50	1
23	28	23.7	81	105	20	0
24	42	32.0	86	120	45	1
25	45	31.5	85	135	50	0
26	54	33.0	92	145	65	1
27	36	29.3	87	115	40	0
28	58	34.0	90	160	60	1
29	29	25.1	84	125	35	0
30	62	36.8	93	150	65	1

31	44	28.6	82	105	30	0
32	37	27.2	85	115	40	1
33	53	32.3	91	130	50	1
34	48	29.4	88	125	45	0
35	40	31.1	85	140	50	1
36	49	30.8	90	145	60	0
37	31	28.0	86	120	40	0
38	50	33.4	92	135	55	1
39	41	29.7	83	130	45	0
40	55	34.3	89	150	65	1
41	46	30.9	88	140	60	0
42	59	35.5	91	160	70	1
43	43	32.8	86	125	50	0
44	32	27.6	87	110	40	0
45	47	33.2	92	145	60	1
46	41	30.0	84	130	55	0
47	60	36.4	90	155	65	1

```
dataset.head()
```

	Age	BMI	Blood Pressure	Glucose Level	Insulin	Outcome
0	25	22.0	80	95	15	0
1	32	27.8	85	120	40	0
2	47	30.5	90	150	60	1
3	55	31.2	88	110	55	1
4	29	24.7	78	100	20	0

```
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48 entries, 0 to 47
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   48 non-null    int64
1   BMI                   48 non-null    float64
2   Blood Pressure        48 non-null    int64
3   Glucose Level         48 non-null    int64
4   Insulin               48 non-null    int64
5   Outcome               48 non-null    int64
dtypes: float64(1), int64(5)
memory usage: 2.4 KB
```

```
dataset.describe()
```

	Age	BMI	Blood Pressure	Glucose Level	Insulin
\count	48.000000	48.000000	48.000000	48.000000	48.000000
mean	45.270833	30.289583	87.291667	131.729167	48.750000

std	11.142672	3.728455	4.135772	20.319511	14.458304
min	25.000000	22.000000	78.000000	95.000000	15.000000
25%	37.000000	27.950000	85.000000	115.000000	40.000000
50%	44.500000	30.450000	88.000000	130.000000	50.000000
75%	54.250000	33.125000	90.250000	146.250000	60.000000
max	70.000000	36.800000	95.000000	180.000000	70.000000

	Outcome
count	48.000000
mean	0.500000
std	0.505291
min	0.000000
25%	0.000000
50%	0.500000
75%	1.000000
max	1.000000

```
X = dataset.iloc[:, [2, 3]].values
y = dataset.iloc[:, 4].values
```

```
print(X[:3, :])
print('-'*15)
print(y[:3])
```

```
[[ 80  95]
 [ 85 120]
 [ 90 150]]
-----
[15 40 60]
```

```
dataset.tail()
```

	Age	BMI	Blood Pressure	Glucose Level	Insulin	Outcome
43	43	32.8	86	125	50	0
44	32	27.6	87	110	40	0
45	47	33.2	92	145	60	1
46	41	30.0	84	130	55	0
47	60	36.4	90	155	65	1

```
dataset["Outcome"].value_counts(normalize=True)
```

```
Outcome
0    0.5
1    0.5
Name: proportion, dtype: float64
```

```
x=dataset.drop(["Outcome"],axis=1)
```

```
y=dataset["Outcome"]
```

x

	Age	BMI	Blood Pressure	Glucose Level	Insulin
0	25	22.0	80	95	15
1	32	27.8	85	120	40
2	47	30.5	90	150	60
3	55	31.2	88	110	55
4	29	24.7	78	100	20
5	60	35.0	95	180	70
6	40	28.4	82	125	45
7	70	34.6	92	160	65
8	50	29.0	89	140	50
9	38	23.6	85	105	25
10	41	33.1	87	130	50
11	45	26.3	90	115	35
12	51	32.0	91	145	55
13	37	28.5	80	105	30
14	60	36.0	89	160	65
15	30	24.3	78	98	20
16	63	34.1	92	155	60
17	35	25.5	86	110	40
18	56	29.5	93	150	60
19	43	28.2	88	120	45
20	64	35.6	91	165	70
21	33	27.0	84	130	55
22	39	30.4	90	140	50
23	28	23.7	81	105	20
24	42	32.0	86	120	45
25	45	31.5	85	135	50
26	54	33.0	92	145	65
27	36	29.3	87	115	40
28	58	34.0	90	160	60
29	29	25.1	84	125	35
30	62	36.8	93	150	65
31	44	28.6	82	105	30
32	37	27.2	85	115	40
33	53	32.3	91	130	50
34	48	29.4	88	125	45
35	40	31.1	85	140	50
36	49	30.8	90	145	60
37	31	28.0	86	120	40
38	50	33.4	92	135	55
39	41	29.7	83	130	45
40	55	34.3	89	150	65
41	46	30.9	88	140	60
42	59	35.5	91	160	70

43	43	32.8	86	125	50
44	32	27.6	87	110	40
45	47	33.2	92	145	60
46	41	30.0	84	130	55
47	60	36.4	90	155	65

y

0	0
1	0
2	1
3	1
4	0
5	1
6	0
7	1
8	1
9	0
10	1
11	0
12	1
13	0
14	1
15	0
16	1
17	0
18	1
19	0
20	1
21	0
22	1
23	0
24	1
25	0
26	1
27	0
28	1
29	0
30	1
31	0
32	1
33	1
34	0
35	1
36	0
37	0
38	1
39	0
40	1
41	0

```
42     1
43     0
44     0
45     1
46     0
47     1
Name: Outcome, dtype: int64
```

```
from sklearn.model_selection import train_test_split
```

```
train_x, test_x, train_y , test_y
=train_test_split(x,y,test_size=0.2,random_state=42)
```

```
train_x
```

	Age	BMI	Blood Pressure	Glucose Level	Insulin
8	50	29.0	89	140	50
3	55	31.2	88	110	55
6	40	28.4	82	125	45
39	41	29.7	83	130	45
33	53	32.3	91	130	50
13	37	28.5	80	105	30
17	35	25.5	86	110	40
45	47	33.2	92	145	60
15	30	24.3	78	98	20
9	38	23.6	85	105	25
16	63	34.1	92	155	60
29	29	25.1	84	125	35
32	37	27.2	85	115	40
46	41	30.0	84	130	55
0	25	22.0	80	95	15
31	44	28.6	82	105	30
30	62	36.8	93	150	65
5	60	35.0	95	180	70
11	45	26.3	90	115	35
34	48	29.4	88	125	45
1	32	27.8	85	120	40
44	32	27.6	87	110	40
21	33	27.0	84	130	55
2	47	30.5	90	150	60
36	49	30.8	90	145	60
35	40	31.1	85	140	50
23	28	23.7	81	105	20
41	46	30.9	88	140	60
10	41	33.1	87	130	50
22	39	30.4	90	140	50
18	56	29.5	93	150	60
47	60	36.4	90	155	65
20	64	35.6	91	165	70
7	70	34.6	92	160	65

42	59	35.5	91	160	70
14	60	36.0	89	160	65
28	58	34.0	90	160	60
38	50	33.4	92	135	55

test_x

	Age	BMI	Blood Pressure	Glucose Level	Insulin
27	36	29.3	87	115	40
40	55	34.3	89	150	65
26	54	33.0	92	145	65
43	43	32.8	86	125	50
24	42	32.0	86	120	45
37	31	28.0	86	120	40
12	51	32.0	91	145	55
19	43	28.2	88	120	45
4	29	24.7	78	100	20
25	45	31.5	85	135	50

train_y

8	1
3	1
6	0
39	0
33	1
13	0
17	0
45	1
15	0
9	0
16	1
29	0
32	1
46	0
0	0
31	0
30	1
5	1
11	0
34	0
1	0
44	0
21	0
2	1
36	0
35	1
23	0
41	0
10	1

```
22    1
18    1
47    1
20    1
7     1
42    1
14    1
28    1
38    1
```

```
Name: Outcome, dtype: int64
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler=MinMaxScaler()
scaler
```

```
MinMaxScaler()
```

```
cols=train_x.columns
cols
```

```
Index(['Age', 'BMI', 'Blood Pressure', 'Glucose Level', 'Insulin'],
      dtype='object')
```

```
train_x_scaled=scaler.fit_transform(train_x)
train_x_scaled
```

```
array([[0.55555556, 0.47297297, 0.64705882, 0.52941176, 0.63636364],
       [0.66666667, 0.62162162, 0.58823529, 0.17647059, 0.72727273],
       [0.33333333, 0.43243243, 0.23529412, 0.35294118, 0.54545455],
       [0.35555556, 0.52027027, 0.29411765, 0.41176471, 0.54545455],
       [0.62222222, 0.69594595, 0.76470588, 0.41176471, 0.63636364],
       [0.26666667, 0.43918919, 0.11764706, 0.11764706, 0.27272727],
       [0.22222222, 0.23648649, 0.47058824, 0.17647059, 0.45454545],
       [0.48888889, 0.75675676, 0.82352941, 0.58823529, 0.81818182],
       [0.11111111, 0.15540541, 0., 0.03529412, 0.09090909],
       [0.28888889, 0.10810811, 0.41176471, 0.11764706, 0.18181818],
       [0.84444444, 0.81756757, 0.82352941, 0.70588235, 0.81818182],
       [0.08888889, 0.20945946, 0.35294118, 0.35294118, 0.36363636],
       [0.26666667, 0.35135135, 0.41176471, 0.23529412, 0.45454545],
       [0.35555556, 0.54054054, 0.35294118, 0.41176471, 0.72727273],
       [0., 0., 0.11764706, 0., 0.],
       [0.42222222, 0.44594595, 0.23529412, 0.11764706, 0.27272727],
       [0.82222222, 1., 0.88235294, 0.64705882, 0.90909091],
       [0.77777778, 0.87837838, 1., 1., 1.],
       [0.44444444, 0.29054054, 0.70588235, 0.23529412, 0.36363636],
       [0.51111111, 0.5, 0.58823529, 0.35294118, 0.54545455],
       [0.15555556, 0.39189189, 0.41176471, 0.29411765, 0.45454545],
       [0.15555556, 0.37837838, 0.52941176, 0.17647059, 0.45454545],
       [0.17777778, 0.33783784, 0.35294118, 0.41176471, 0.72727273],
       [0.48888889, 0.57432432, 0.70588235, 0.64705882, 0.81818182],
```



```
[0.53333333, 0.59459459, 0.70588235, 0.58823529, 0.81818182],
[0.33333333, 0.61486486, 0.41176471, 0.52941176, 0.63636364],
[0.06666667, 0.11486486, 0.17647059, 0.11764706, 0.09090909],
[0.46666667, 0.60135135, 0.58823529, 0.52941176, 0.81818182],
[0.35555556, 0.75, 0.52941176, 0.41176471, 0.63636364],
[0.31111111, 0.56756757, 0.70588235, 0.52941176, 0.63636364],
[0.68888889, 0.50675676, 0.88235294, 0.64705882, 0.81818182],
[0.77777778, 0.97297297, 0.70588235, 0.70588235, 0.90909091],
[0.86666667, 0.91891892, 0.76470588, 0.82352941, 1.],
[1., 0.85135135, 0.82352941, 0.76470588, 0.90909091],
[0.75555556, 0.91216216, 0.76470588, 0.76470588, 1.],
[0.77777778, 0.94594595, 0.64705882, 0.76470588, 0.90909091],
[0.73333333, 0.81081081, 0.70588235, 0.76470588, 0.81818182],
[0.55555556, 0.77027027, 0.82352941, 0.47058824, 0.72727273]])
```

```
train_x_scaled=pd.DataFrame(train_x_scaled,columns=cols)
```

```
train_x_scaled
```

	Age	BMI	Blood Pressure	Glucose Level	Insulin
0	0.555556	0.472973	0.647059	0.529412	0.636364
1	0.666667	0.621622	0.588235	0.176471	0.727273
2	0.333333	0.432432	0.235294	0.352941	0.545455
3	0.355556	0.520270	0.294118	0.411765	0.545455
4	0.622222	0.695946	0.764706	0.411765	0.636364
5	0.266667	0.439189	0.117647	0.117647	0.272727
6	0.222222	0.236486	0.470588	0.176471	0.454545
7	0.488889	0.756757	0.823529	0.588235	0.818182
8	0.111111	0.155405	0.000000	0.035294	0.090909
9	0.288889	0.108108	0.411765	0.117647	0.181818
10	0.844444	0.817568	0.823529	0.705882	0.818182
11	0.088889	0.209459	0.352941	0.352941	0.363636
12	0.266667	0.351351	0.411765	0.235294	0.454545
13	0.355556	0.540541	0.352941	0.411765	0.727273
14	0.000000	0.000000	0.117647	0.000000	0.000000
15	0.422222	0.445946	0.235294	0.117647	0.272727
16	0.822222	1.000000	0.882353	0.647059	0.909091
17	0.777778	0.878378	1.000000	1.000000	1.000000
18	0.444444	0.290541	0.705882	0.235294	0.363636
19	0.511111	0.500000	0.588235	0.352941	0.545455
20	0.155556	0.391892	0.411765	0.294118	0.454545
21	0.155556	0.378378	0.529412	0.176471	0.454545
22	0.177778	0.337838	0.352941	0.411765	0.727273
23	0.488889	0.574324	0.705882	0.647059	0.818182
24	0.533333	0.594595	0.705882	0.588235	0.818182
25	0.333333	0.614865	0.411765	0.529412	0.636364
26	0.066667	0.114865	0.176471	0.117647	0.090909
27	0.466667	0.601351	0.588235	0.529412	0.818182
28	0.355556	0.750000	0.529412	0.411765	0.636364
29	0.311111	0.567568	0.705882	0.529412	0.636364

30	0.688889	0.506757	0.882353	0.647059	0.818182
31	0.777778	0.972973	0.705882	0.705882	0.909091
32	0.866667	0.918919	0.764706	0.823529	1.000000
33	1.000000	0.851351	0.823529	0.764706	0.909091
34	0.755556	0.912162	0.764706	0.764706	1.000000
35	0.777778	0.945946	0.647059	0.764706	0.909091
36	0.733333	0.810811	0.705882	0.764706	0.818182
37	0.555556	0.770270	0.823529	0.470588	0.727273

```

from sklearn.linear_model import LogisticRegression as LogReg
logreg=LogReg()
logreg.fit(train_x,train_y)
LogisticRegression()
train_predict=logreg.predict(train_x)
test_predict=logreg.predict(test_x)

train_predict
array([1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0,
       0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1], dtype=int64)
test_predict
array([0, 1, 1, 1, 1, 0, 1, 0, 0, 1], dtype=int64)

from sklearn.metrics import f1_score, confusion_matrix, roc_auc_score,
roc_curve
f1_score(train_predict,train_y)
0.9047619047619048
f1_score(test_predict,test_y)
0.8
conf1=confusion_matrix(train_y,train_predict)
conf1
array([[15,  3],
       [ 1, 19]], dtype=int64)

from sklearn.metrics import accuracy_score, confusion_matrix
accuracy = accuracy_score(test_y, test_predict)
conf_matrix = confusion_matrix(test_y, test_predict)
accuracy

```

0.8

conf_matrix

```
array([[4, 2],  
       [0, 4]], dtype=int64)
```

```
from sklearn.metrics import classification_report
```

```
print("Accuracy:", accuracy)  
print("Confusion Matrix:")  
print(conf_matrix)  
print("\nClassification Report:")  
print(classification_report(test_y, test_predict))
```

Accuracy: 0.8

Confusion Matrix:

```
[[4 2]  
 [0 4]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.67	0.80	6
1	0.67	1.00	0.80	4
accuracy			0.80	10
macro avg	0.83	0.83	0.80	10
weighted avg	0.87	0.80	0.80	10

```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay  
import matplotlib.pyplot as plt
```


```
# Compute the confusion matrix
```

```
conf_matrix = confusion_matrix(test_y, test_predict)
```

```
# Display the confusion matrix
```

```
disp = ConfusionMatrixDisplay(conf_matrix)
```

```
disp.plot(cmap="Blues") # You can change the color map if needed  
plt.show()
```



```
from sklearn.metrics import confusion_matrix

# Compute confusion matrix
conf_matrix = confusion_matrix(test_y, test_predict)

# Extract values
true_negative = conf_matrix[0][0] # TN
false_positive = conf_matrix[0][1] # FP
false_negative = conf_matrix[1][0] # FN
true_positive = conf_matrix[1][1] # TP

# Print values
print(f"True Negative: {true_negative}")
print(f"False Positive: {false_positive}")
print(f"False Negative: {false_negative}")
print(f"True Positive: {true_positive}")

True Negative: 4
False Positive: 2
False Negative: 0
True Positive: 4

Accuracy = (true_positive + true_negative) / (true_positive
+false_positive)
```

```

Accuracy
# Precision
Precision = true_positive/(true_positive+false_positive)
Precision
# Recall
Recall = true_positive/(true_positive+false_negative)
Recall
# F1 Score
F1_Score = 2*(Recall * Precision) / (Recall + Precision)
F1_Score

0.8

Accuracy
1.3333333333333333

Precision
0.6666666666666666

Recall
1.0

F1_Score
0.8


auc_score=roc_auc_score(test_y,test_predict)

fpr,tpr,thresholds=roc_curve(test_y,test_predict)


thresholds
array([inf,  1.,  0.])

plt.plot(fpr, tpr, color='orange', label='ROC')
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--',label='ROC
curve (area = 0.85)')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()

```



```
import seaborn as sns
sns.heatmap(conf_matrix, annot=True)
<Axes: >
```



Name: Shruti Manwar Roll no :13229